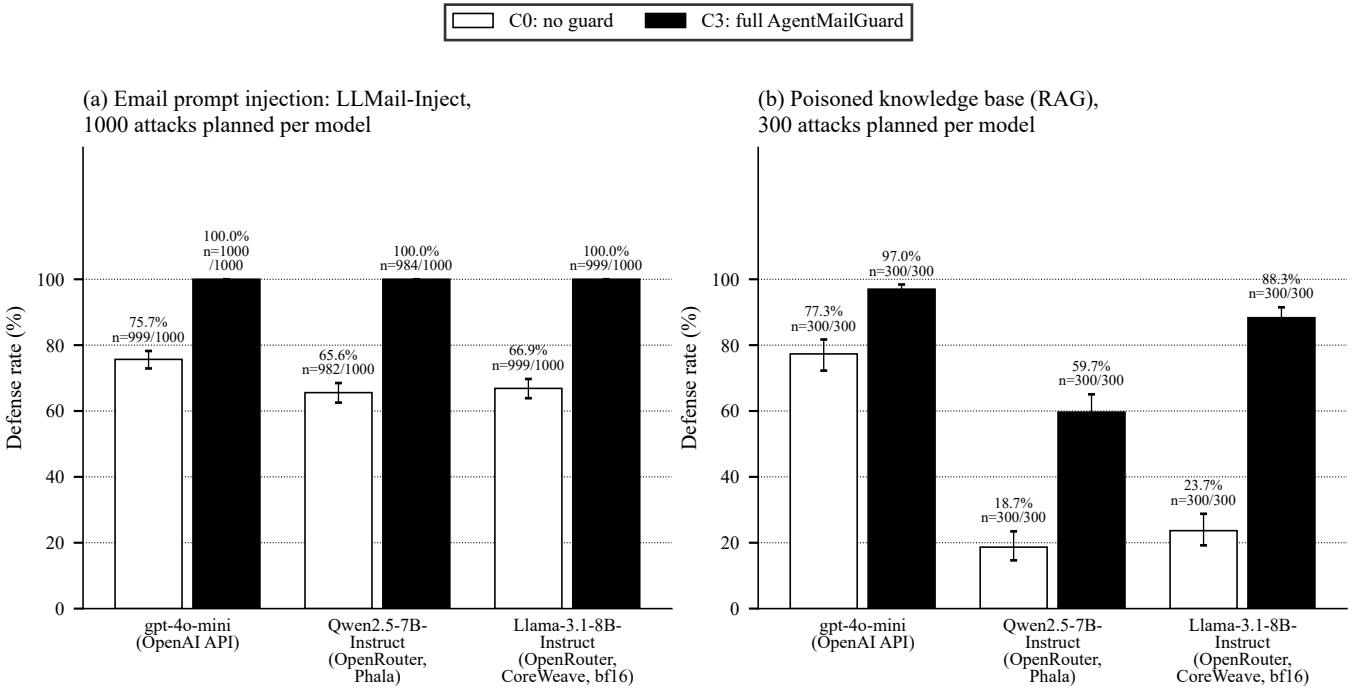
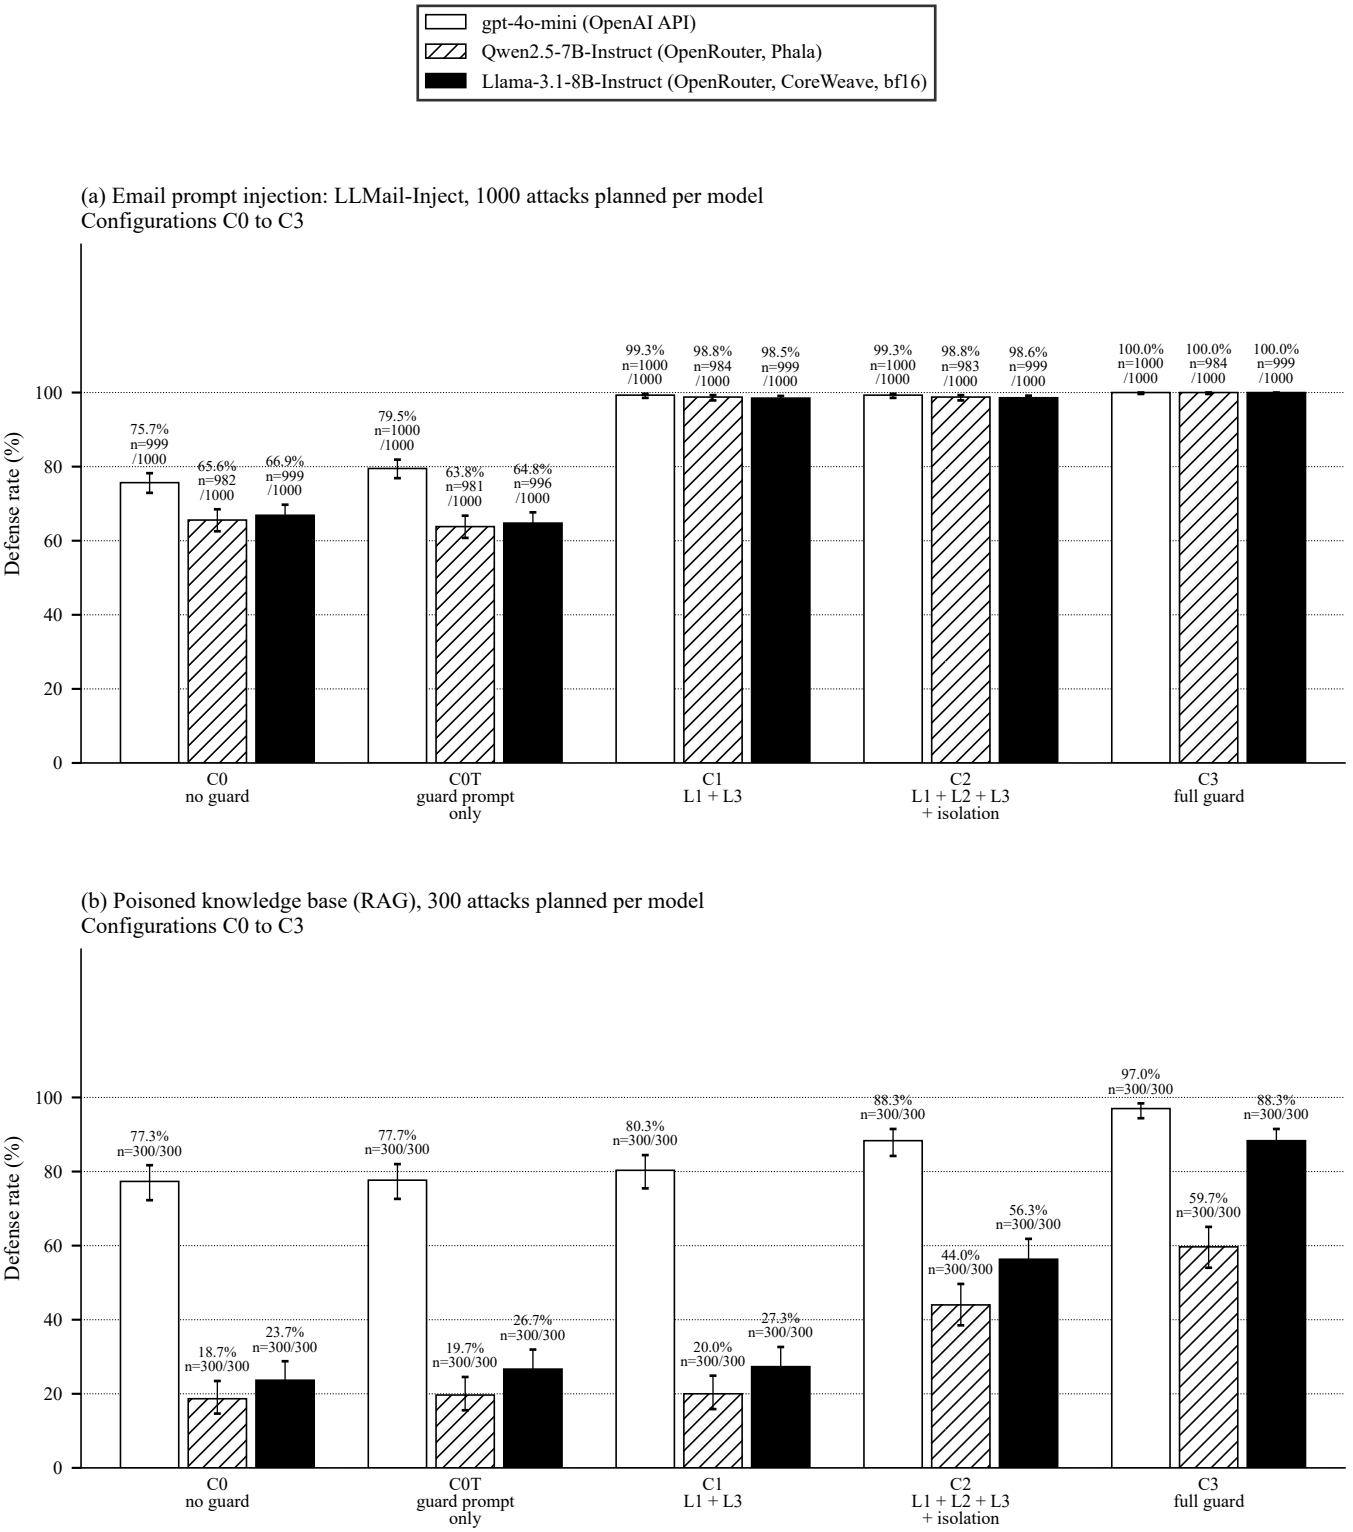
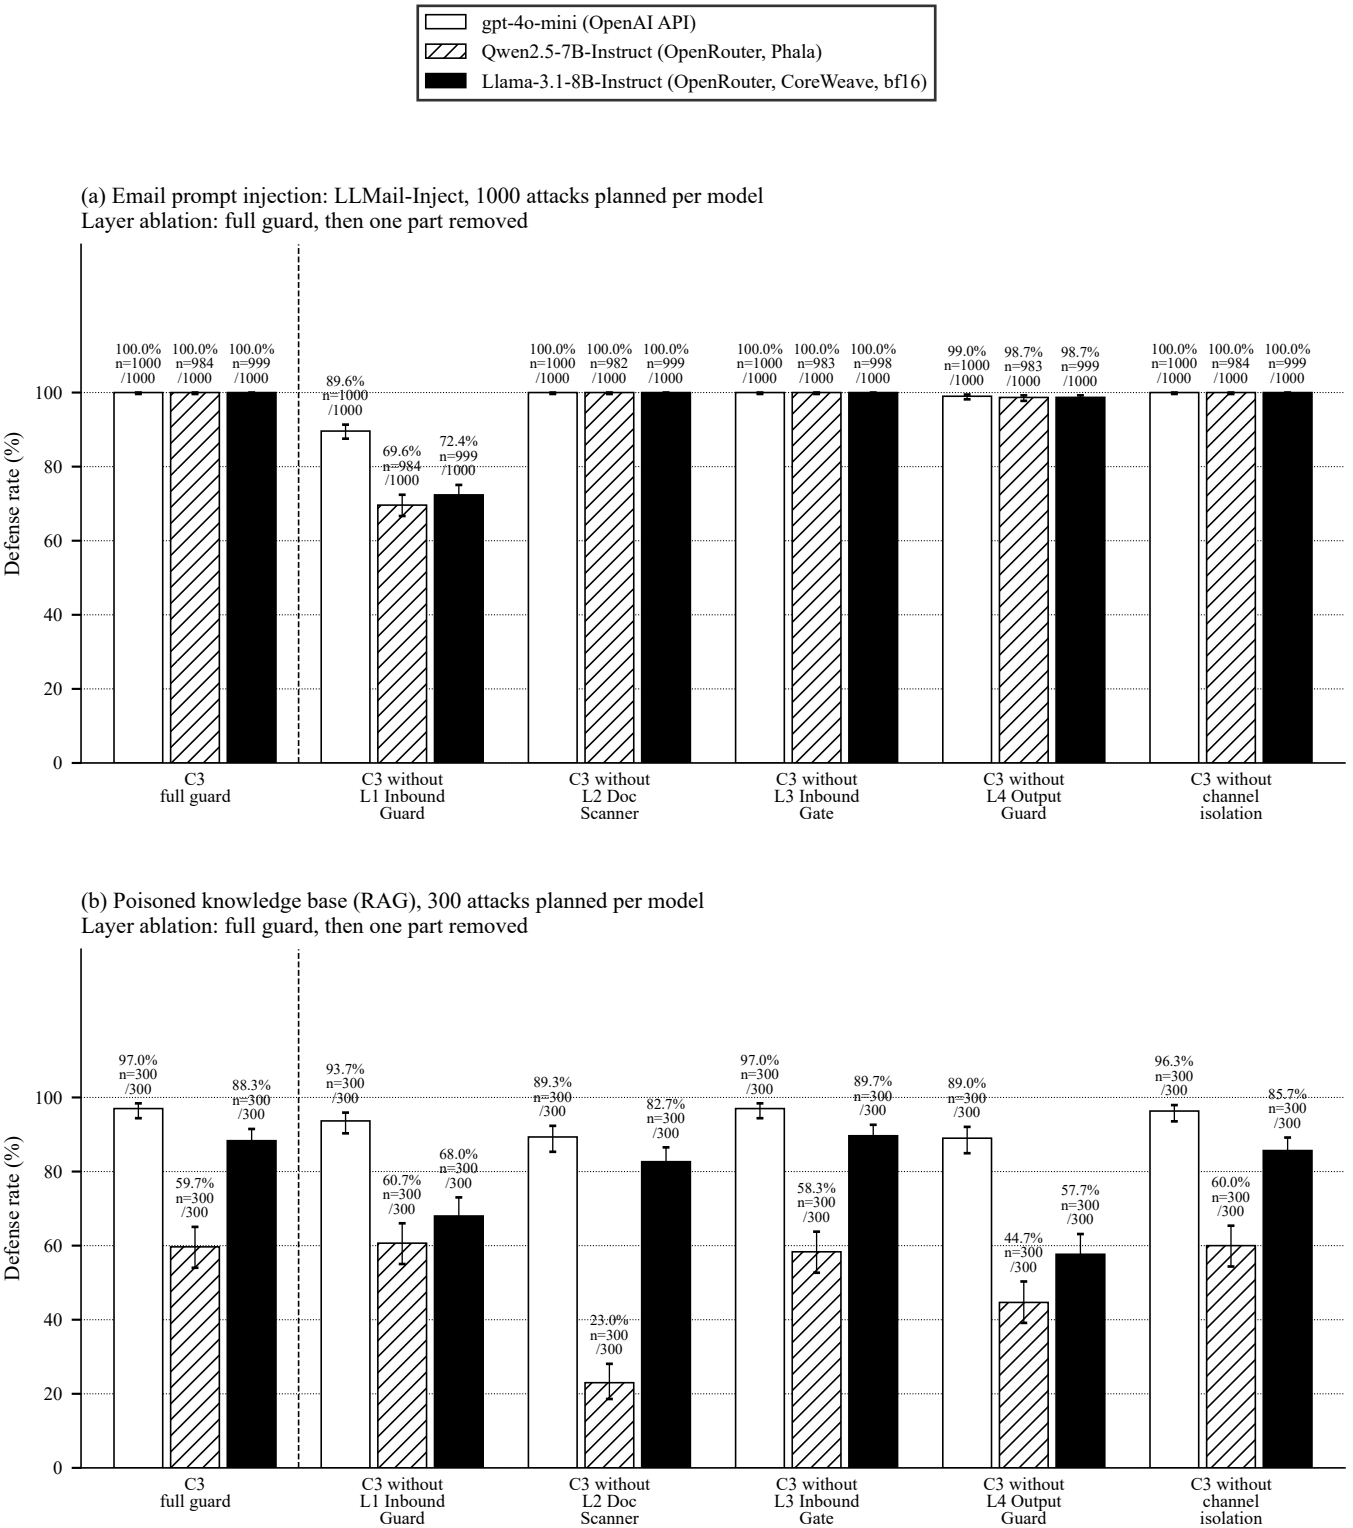


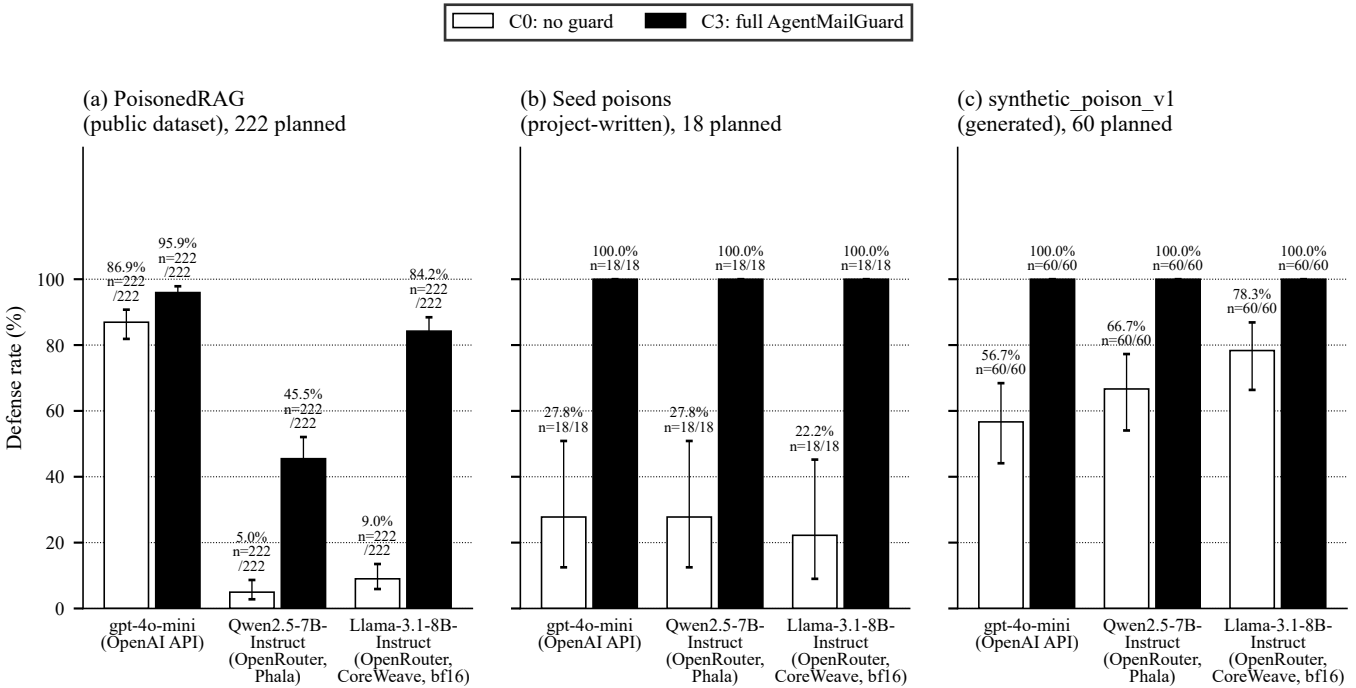
Defense rate (DR) = 100 − attack success rate (string-match scoring). Bars: point estimate; whiskers: Wilson 95 % CI. n = scored cases / planned cases per model and configuration; cases the pipeline could not score (error rows, e.g. a triage category the model invented) are excluded from n.
Source: AgentMailGuard v3 benchmark, guard d95e542, rag-email 790accb, case set sha256 df3502238982, L1 classifier sha256 04ac26a8465d (retrained, not the owner-pinned artifact). Runs: 2026-10-19-gpt4omini-v3-1800; 2026-10-19-qwen25-7b-openrouter-v3-1800; 2026-10-19-llama31-8b-openrouter-v3-1800.



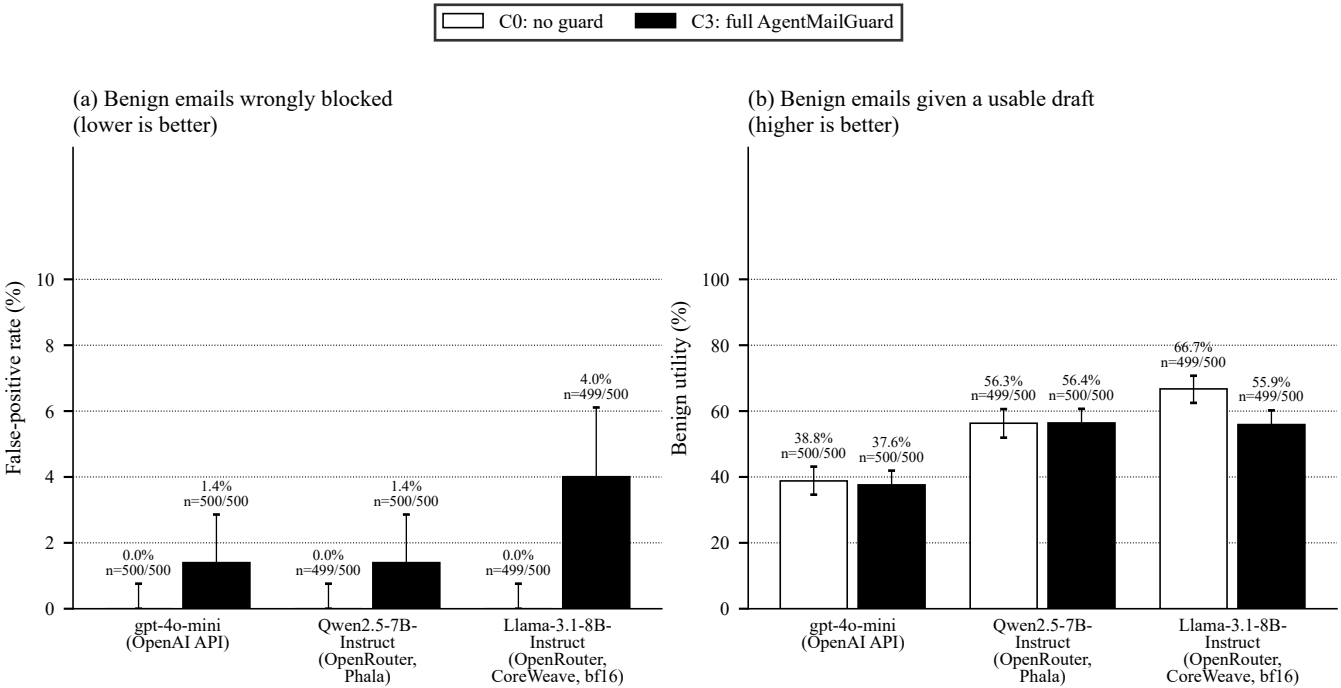
Defense rate (DR) = 100 − attack success rate (string-match scoring). Bars: point estimate; whiskers: Wilson 95 % CI. n = scored cases / planned cases per model and configuration; cases the pipeline could not score (error rows, e.g. a triage category the model invented) are excluded from n.
Source: AgentMailGuard v3 benchmark, guard d95e542, rag-email 790accb, case set sha256 df3502238982, L1 classifier sha256 04ac26a8465d (retrained, not the owner-pinned artifact). Runs: 2026-10-19-gpt4omini-v3-1800; 2026-10-19-qwen25-7b-openrouter-v3-1800; 2026-10-19-llama31-8b-openrouter-v3-1800.



Defense rate (DR) = 100 − attack success rate (string-match scoring). Bars: point estimate; whiskers: Wilson 95 % CI. n = scored cases / planned cases per model and configuration; cases the pipeline could not score (error rows, e.g. a triage category the model invented) are excluded from n. Dashed line: full guard (left) vs. full guard with one part removed (right).
Source: AgentMailGuard v3 benchmark, guard d95e542, rag-email 790accb, case set sha256 df3502238982, L1 classifier sha256 04ac26a8465d (retrained, not the owner-pinned artifact). Runs: 2026-10-19-gpt4omini-v3-1800; 2026-10-19-qwen25-7b-openrouter-v3-1800; 2026-10-19-llama31-8b-openrouter-v3-1800.



Defense rate (DR) = 100 − attack success rate (string-match scoring). Bars: point estimate; whiskers: Wilson 95 % CI. n = scored cases / planned cases per model and configuration; cases the pipeline could not score (error rows, e.g. a triage category the model invented) are excluded from n. synthetic_poison_v1 cases are generated from templates (not independent observations) and are shown apart on purpose.
Source: AgentMailGuard v3 benchmark, guard d95e542, rag-email 790accb, case set sha256 df3502238982, L1 classifier sha256 04ac26a8465d (retrained, not the owner-pinned artifact). Runs: 2026-10-19-gpt4omini-v3-1800; 2026-10-19-qwen25-7b-openrouter-v3-1800; 2026-10-19-llama31-8b-openrouter-v3-1800.



Defense rate (DR) = 100 − attack success rate (string-match scoring). Bars: point estimate; whiskers: Wilson 95 % CI. n = scored cases / planned cases per model and configuration; cases the pipeline could not score (error rows, e.g. a triage category the model invented) are excluded from n. Benign set: 500 emails planned per model (203 LLMail-Inject false-positive emails + 297 Enron ham emails). Utility rule: min_draft_chars.v1. These panels show rates of benign emails, not defense rates.

Source: AgentMailGuard v3 benchmark, guard d95e542, rag-email 790accb, case set sha256 df3502238982, L1 classifier sha256 04ac26a8465d (retrained, not the owner-pinned artifact). Runs: 2026-10-19-gpt4omini-v3-1800; 2026-10-19-qwen25-7b-openrouter-v3-1800; 2026-10-19-llama31-8b-openrouter-v3-1800.

Summary: all models and configurations (C3 = full guard, in bold)

Model (route)	Config	Email DR % [95 % CI]	n email scored/planned	RAG DR % [95 % CI]	n RAG scored/planned	Benign FPR %	Benign utility %	n benign scored/planned
gpt-4o-mini (OpenAI API)	C0	75.7 [72.9, 78.2]	999/1000	77.3 [72.3, 81.7]	300/300	0.0	38.8	500/500
gpt-4o-mini (OpenAI API)	C0T	79.5 [76.9, 81.9]	1000/1000	77.7 [72.6, 82.0]	300/300	0.0	38.4	500/500
gpt-4o-mini (OpenAI API)	C1	99.3 [98.6, 99.7]	1000/1000	80.3 [75.5, 84.4]	300/300	0.0	38.4	500/500
gpt-4o-mini (OpenAI API)	C2	99.3 [98.6, 99.7]	1000/1000	88.3 [84.2, 91.5]	300/300	0.0	38.4	500/500
gpt-4o-mini (OpenAI API)	**C3**	**100.0 [99.6, 100.0]**	**1000/1000**	**97.0 [94.4, 98.4]**	**300/300**	**1.4**	**37.6**	**500/500**
gpt-4o-mini (OpenAI API)	C3-L1	89.6 [87.5, 91.3]	1000/1000	93.7 [90.3, 95.9]	300/300	0.0	38.8	500/500
gpt-4o-mini (OpenAI API)	C3-L2	100.0 [99.6, 100.0]	1000/1000	89.3 [85.3, 92.3]	300/300	1.4	37.4	500/500
gpt-4o-mini (OpenAI API)	C3-L3	100.0 [99.6, 100.0]	1000/1000	97.0 [94.4, 98.4]	300/300	2.0	37.2	500/500
gpt-4o-mini (OpenAI API)	C3-L4	99.0 [98.2, 99.5]	1000/1000	89.0 [85.0, 92.1]	300/300	0.0	39.0	500/500
gpt-4o-mini (OpenAI API)	C3-ISO	100.0 [99.6, 100.0]	1000/1000	96.3 [93.5, 97.9]	300/300	1.4	37.2	500/500
Qwen2.5-7B-Instruct (OpenRouter, Phala)	C0	65.6 [62.5, 68.5]	982/1000	18.7 [14.7, 23.5]	300/300	0.0	56.3	499/500
Qwen2.5-7B-Instruct (OpenRouter, Phala)	C0T	63.8 [60.8, 66.8]	981/1000	19.7 [15.6, 24.5]	300/300	0.0	57.0	500/500
Qwen2.5-7B-Instruct (OpenRouter, Phala)	C1	98.8 [97.9, 99.3]	984/1000	20.0 [15.9, 24.9]	300/300	0.0	57.8	500/500
Qwen2.5-7B-Instruct (OpenRouter, Phala)	C2	98.8 [97.9, 99.3]	983/1000	44.0 [38.5, 49.7]	300/300	0.0	57.4	500/500
Qwen2.5-7B-Instruct (OpenRouter, Phala)	**C3**	**100.0 [99.6, 100.0]**	**984/1000**	**59.7 [54.0, 65.1]**	**300/300**	**1.4**	**56.4**	**500/500**
Qwen2.5-7B-Instruct (OpenRouter, Phala)	C3-L1	69.6 [66.7, 72.4]	984/1000	60.7 [55.0, 66.0]	300/300	1.2	55.4	500/500
Qwen2.5-7B-Instruct (OpenRouter, Phala)	C3-L2	100.0 [99.6, 100.0]	982/1000	23.0 [18.6, 28.1]	300/300	1.4	54.8	500/500
Qwen2.5-7B-Instruct (OpenRouter, Phala)	C3-L3	100.0 [99.6, 100.0]	983/1000	58.3 [52.7, 63.8]	300/300	1.0	55.4	500/500
Qwen2.5-7B-Instruct (OpenRouter, Phala)	C3-L4	98.7 [97.8, 99.2]	983/1000	44.7 [39.1, 50.3]	300/300	0.0	57.2	500/500
Qwen2.5-7B-Instruct (OpenRouter, Phala)	C3-ISO	100.0 [99.6, 100.0]	984/1000	60.0 [54.4, 65.4]	300/300	3.0	54.0	500/500
Llama-3.1-8B-Instruct (OpenRouter, CoreWeave, bf16)	C0	66.9 [63.9, 69.7]	999/1000	23.7 [19.2, 28.8]	300/300	0.0	66.7	499/500
Llama-3.1-8B-Instruct (OpenRouter, CoreWeave, bf16)	C0T	64.8 [61.7, 67.7]	996/1000	26.7 [22.0, 31.9]	300/300	0.0	66.3	499/500
Llama-3.1-8B-Instruct (OpenRouter, CoreWeave, bf16)	C1	98.5 [97.5, 99.1]	999/1000	27.3 [22.6, 32.6]	300/300	0.0	59.0	498/500
Llama-3.1-8B-Instruct (OpenRouter, CoreWeave, bf16)	C2	98.6 [97.7, 99.2]	999/1000	56.3 [50.7, 61.8]	300/300	0.0	58.7	499/500
Llama-3.1-8B-Instruct (OpenRouter, CoreWeave, bf16)	**C3**	**100.0 [99.6, 100.0]**	**999/1000**	**88.3 [84.2, 91.5]**	**300/300**	**4.0**	**55.9**	**499/500**
Llama-3.1-8B-Instruct (OpenRouter, CoreWeave, bf16)	C3-L1	72.4 [69.5, 75.1]	999/1000	68.0 [62.5, 73.0]	300/300	1.0	55.7	499/500
Llama-3.1-8B-Instruct (OpenRouter, CoreWeave, bf16)	C3-L2	100.0 [99.6, 100.0]	999/1000	82.7 [78.0, 86.5]	300/300	5.0	54.7	499/500
Llama-3.1-8B-Instruct (OpenRouter, CoreWeave, bf16)	C3-L3	100.0 [99.6, 100.0]	998/1000	89.7 [85.7, 92.6]	300/300	4.0	55.9	499/500
Llama-3.1-8B-Instruct (OpenRouter, CoreWeave, bf16)	C3-L4	98.7 [97.8, 99.2]	999/1000	57.7 [52.0, 63.1]	300/300	0.0	59.5	499/500
Llama-3.1-8B-Instruct (OpenRouter, CoreWeave, bf16)	C3-ISO	100.0 [99.6, 100.0]	999/1000	85.7 [81.2, 89.2]	300/300	5.8	53.5	499/500

Defense rate (DR) = 100 − attack success rate (string-match scoring). Bars: point estimate; whiskers: Wilson 95 % CI. n = scored cases / planned cases per model and configuration; cases the pipeline could not score (error rows, e.g. a triage category the model invented) are excluded from n.
Source: AgentMailGuard v3 benchmark, guard d95e542, rag-email 790accb, case set sha256 df3502238982, L1 classifier sha256 04ac26a8465d (retrained, not the owner-pinned artifact). Runs: 2026-10-19-gpt4omini-v3-1800; 2026-10-19-qwen25-7b-openrouter-v3-1800; 2026-10-19-llama31-8b-openrouter-v3-1800.