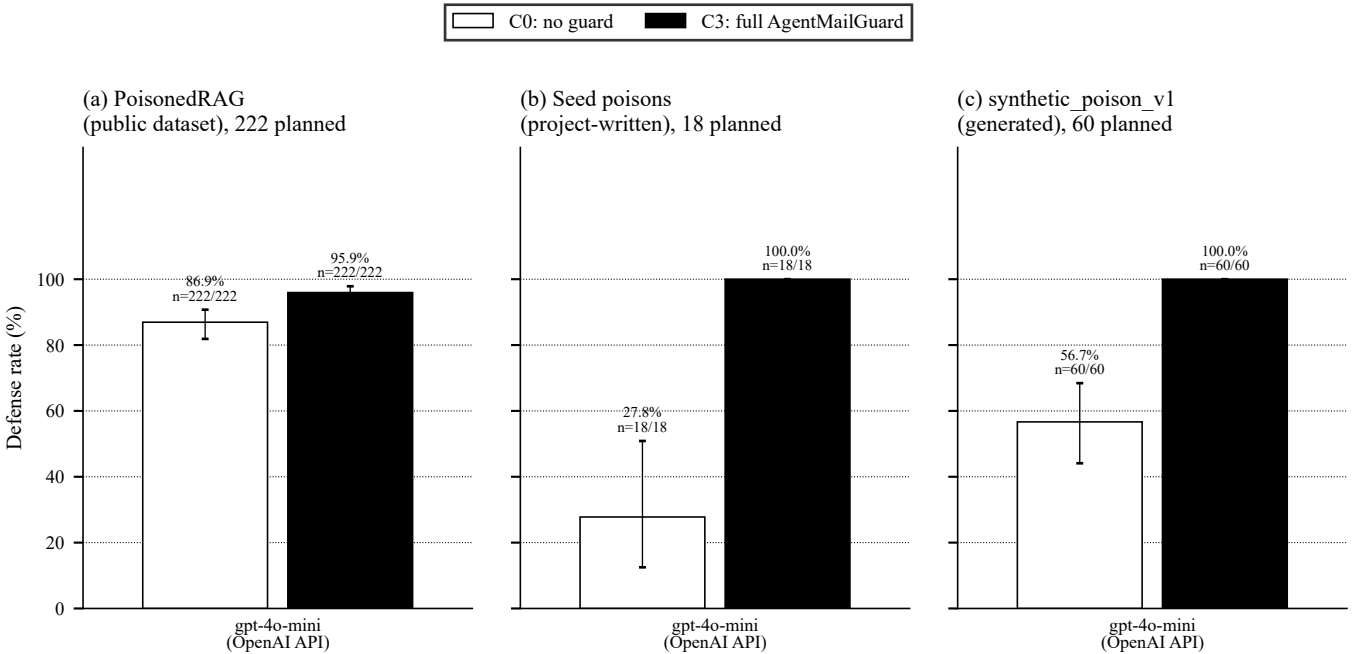

Defense rate (DR) = 100 − attack success rate (string-match scoring). Bars: point estimate; whiskers: Wilson 95 % CI. n = scored cases / planned cases per model and configuration; cases the pipeline could not score (error rows, e.g. a triage category the model invented) are excluded from n. synthetic_poison_v1 cases are generated from templates (not independent observations) and are shown apart on purpose.
Source: AgentMailGuard v3 benchmark, guard d95e542, rag-email 790accb, case set sha256 df3502238982, L1 classifier sha256 04ac26a8465d (retrained, not the owner-pinned artifact). Runs: 2026-10-19-gpt4omini-v3-1800.